

```

from sklearn.datasets import fetch_openml
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPClassifier
from sklearn.metrics import classification_report, confusion_matrix
import matplotlib.pyplot as plt

# Load the MNIST dataset
mnist = fetch_openml('mnist_784', version=1)

# Features and labels
X, y = mnist['data'], mnist['target']

# Convert labels to integers
y = y.astype(int)

# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the data
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

# Initialize the MLPClassifier
mlp = MLPClassifier(hidden_layer_sizes=(64, 64), max_iter=20, alpha=1e-4,
                    solver='sgd', verbose=10, random_state=1,
                    learning_rate_init=.1)

# Train the model
mlp.fit(X_train, y_train)

↩ Iteration 1, loss = 0.28395107
Iteration 2, loss = 0.15718102
Iteration 3, loss = 0.13286050
Iteration 4, loss = 0.10744293
Iteration 5, loss = 0.09562792
Iteration 6, loss = 0.08628949
Iteration 7, loss = 0.08982760
Iteration 8, loss = 0.07116525
Iteration 9, loss = 0.06793491
Iteration 10, loss = 0.07117720
Iteration 11, loss = 0.07209220
Iteration 12, loss = 0.06278195
Iteration 13, loss = 0.05612610
Iteration 14, loss = 0.06040064
Iteration 15, loss = 0.06829053
Iteration 16, loss = 0.06216053
Iteration 17, loss = 0.09009062
Iteration 18, loss = 0.16276313
Iteration 19, loss = 0.09797132
Iteration 20, loss = 0.08373181
/usr/local/lib/python3.10/dist-packages/sklearn/neural_network/_multilayer_perceptron.py:690: ConvergenceWarning: Stochastic Optimizer:
warnings.warn(
  MLPClassifier
  MLPClassifier(hidden_layer_sizes=(64, 64), learning_rate_init=0.1, max_iter=20,
                random_state=1, solver='sgd', verbose=10)

# Predict the labels for the test set
y_pred = mlp.predict(X_test)

# Display classification report
print("Classification Report:\n", classification_report(y_test, y_pred))

# Display confusion matrix
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

```

↩ Classification Report:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.97	0.98	0.97	1343
1	0.97	0.98	0.98	1600
2	0.89	0.96	0.92	1380
3	0.98	0.94	0.96	1433
4	0.96	0.97	0.96	1295
5	0.95	0.94	0.94	1273
6	0.95	0.96	0.95	1396
7	0.96	0.96	0.96	1503
8	0.98	0.90	0.94	1357
9	0.94	0.94	0.94	1420
accuracy			0.95	14000
macro avg	0.95	0.95	0.95	14000
weighted avg	0.95	0.95	0.95	14000

Confusion Matrix:

```
[[1314  2  9  0  0  1 11  3  2  1]
 [  1 1570  8  4  3  3  3  6  2  0]
 [  3  3 1325  4  7  5 10 11  8  4]
 [  0  1  32 1342  1 29  2 13  2 11]
 [  2  3 12  0 1250  1  6  3  2 16]
 [  3  5 11  5  3 1199 35  1  4  7]
 [ 16  1 24  0  7  3 1345  0  0  0]
 [  4  2 17  3  6  4  0 1448  1 18]
 [ 10 18 30  9  3 15 10  9 1223 30]
 [  5  8 18  8 22  6  1 18  5 1329]]
```

# Function to display an image from the MNIST dataset

def plot\_sample(X, y, index):

plt.imshow(X[index].reshape(28, 28), cmap='gray')

plt.title(f"True Label: {y[index]}")

plt.show()


# Display a few test samples with predicted labels

for i in range(5):


plot\_sample(X\_test, y\_pred, i)




True Label: 8



True Label: 4




True Label: 8

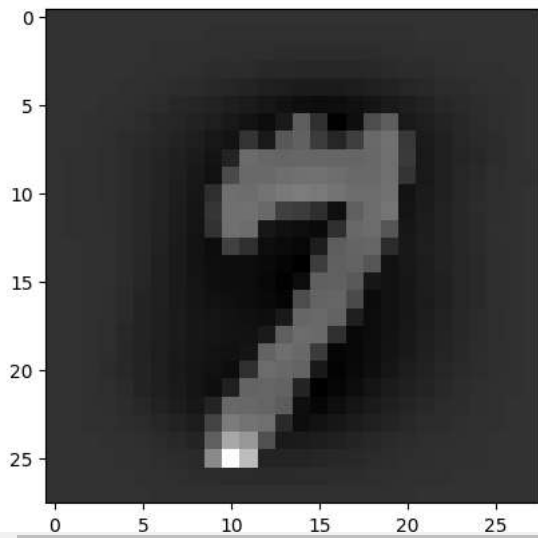


True Label: 7





True Label: 7



```
import matplotlib.pyplot as plt
# Plot the loss curve during training
plt.plot(mlp.loss_curve_)
plt.title("Loss Curve during Training")
plt.xlabel("Epochs")
plt.ylabel("Loss")
plt.show()
```

