

```
import numpy as np
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import LSTM, Dense, Embedding, SpatialDropout1D
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.datasets import imdb
from sklearn.metrics import accuracy_score
```

```
# Load the IMDb dataset (keep only the top 10,000 most frequent words)
max_words = 10000
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_words)
```

```
# Check sample data
print(f"Sample Review (Token IDs): {x_train[0]}")
print(f"Sentiment: {y_train[0]}") # 1 = positive, 0 = negative
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz>  
 17464789/17464789 — 0s 0us/step  
 Sample Review (Token IDs): [1, 14, 22, 16, 43, 530, 973, 1622, 1385, 65, 458, 4468, 66, 3941, 4, 173, 36, 256, 5, 25, 100, 43, 838, 112,  
 Sentiment: 1

```
max_sequence_length = 100 # Maximum review length
```

```
# Pad the sequences to ensure uniform input size
x_train = pad_sequences(x_train, maxlen=max_sequence_length)
x_test = pad_sequences(x_test, maxlen=max_sequence_length)
```

```
print(f"Padded Review: {x_train[0]}")
print(f"Shape of Training Data: {x_train.shape}")
```

Padded Review: [1415 33 6 22 12 215 28 77 52 5 14 407 16 82  
 2 8 4 107 117 5952 15 256 4 2 7 3766 5 723  
 36 71 43 530 476 26 400 317 46 7 4 2 1029 13  
 104 88 4 381 15 297 98 32 2071 56 26 141 6 194  
 7486 18 4 226 22 21 134 476 26 480 5 144 30 5535  
 18 51 36 28 224 92 25 104 4 226 65 16 38 1334  
 88 12 16 283 5 16 4472 113 103 32 15 16 5345 19  
 178 32]  
 Shape of Training Data: (25000, 100)

```
embedding_dim = 100 # Dimension of the word embeddings
```

```
model = Sequential()
model.add(Embedding(input_dim=max_words, output_dim=embedding_dim, input_length=max_sequence_length))
model.add(SpatialDropout1D(0.2)) # Dropout to prevent overfitting
model.add(LSTM(100, dropout=0.2, recurrent_dropout=0.2))
model.add(Dense(1, activation='sigmoid')) # Binary output (positive/negative)
```

```
# Compile the model
model.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
# Model summary
model.summary()
```

/usr/local/lib/python3.10/dist-packages/keras/src/layers/core/embedding.py:90: UserWarning: Argument `input\_length` is deprecated. Just  
 warnings.warn(  
 Model: "sequential"

Layer (type)	Output Shape	Param #
embedding ( <a href="#">Embedding</a> )	?	0 (unbuilt)
spatial_dropout1d ( <a href="#">SpatialDropout1D</a> )	?	0 (unbuilt)
lstm ( <a href="#">LSTM</a> )	?	0 (unbuilt)
dense ( <a href="#">Dense</a> )	?	0 (unbuilt)

Total params: 0 (0.00 B)  
 Trainable params: 0 (0.00 B)

```
batch_size = 64
epochs = 3
```

```
history = model.fit(x_train, y_train, validation_data=(x_test, y_test),
                    batch_size=batch_size, epochs=epochs, verbose=1)
```

```
Epoch 1/3
391/391 ————— 105s 262ms/step - accuracy: 0.7009 - loss: 0.5526 - val_accuracy: 0.8310 - val_loss: 0.3862
Epoch 2/3
391/391 ————— 106s 270ms/step - accuracy: 0.8657 - loss: 0.3266 - val_accuracy: 0.8466 - val_loss: 0.3583
Epoch 3/3
391/391 ————— 104s 265ms/step - accuracy: 0.8909 - loss: 0.2777 - val_accuracy: 0.8459 - val_loss: 0.3695
```

```
# Evaluate the model on test data
test_loss, test_accuracy = model.evaluate(x_test, y_test, verbose=0)
print(f"Test Accuracy: {test_accuracy:.2f}")
```

```
# Make predictions
y_pred = (model.predict(x_test) > 0.5).astype("int32")
print(f"Sample Prediction: {y_pred[0]}")
```

```
Test Accuracy: 0.85
782/782 ————— 22s 28ms/step
Sample Prediction: [0]
```

```
import matplotlib.pyplot as plt
```

```
# Plot training & validation accuracy
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.show()
```



